

LAB 8

Aim: To implement Mutual Exclusion / Deadlock Detection

Lab Outcome:

Demonstrate Mutual Exclusion algorithms and deadlock handling

Theory:

Mutual Exclusion in Distributed System:

Mutual exclusion is a concurrency control property which is introduced to prevent race conditions. It is the requirement that a process cannot enter its critical section while another concurrent process is currently present or executing in its critical section i.e., only one process is allowed to execute the critical section at any given instance of time.

In Distributed systems, we neither have shared memory nor a common physical clock and therefore we cannot solve mutual exclusion problem using shared variables. To eliminate the mutual exclusion problem in distributed system approach based on message passing is used.

Requirements of Mutual exclusion Algorithm:

- **No Deadlock:**
Two or more site should not endlessly wait for any message that will never arrive.
- **No Starvation:**
Any site should not wait indefinitely to execute critical section while other sites are repeatedly executing critical section
- **Fairness:**
Each site should get a fair chance to execute critical section. Any request to execute critical section must be executed in the order they are made.
- **Fault Tolerance:**
In case of failure, it should be able to recognize it by itself in order to continue functioning without any disruption.

Solution to distributed mutual exclusion:

As we know shared variables or a local kernel cannot be used to implement mutual exclusion in distributed systems. Message passing is a way to implement mutual exclusion. Below are the three approaches based on message passing to implement mutual exclusion in distributed systems:

Token Based Algorithm:

- A unique token is shared among all the sites.
- If a site possesses the unique token, it is allowed to enter its critical section
- This approach uses sequence number to order requests for the critical section.
- Each requests for critical section contains a sequence number. This sequence number is used to distinguish old and current requests.
- This approach insures Mutual exclusion as the token is unique
- Example: Suzuki-Kasami's Broadcast Algorithm

Non-token based approach:

- A site communicates with other sites in order to determine which sites should execute critical section next. This requires exchange of two or more successive round of messages among sites.
- This approach use timestamps instead of sequence number to order requests for the critical section.
- When ever a site make request for critical section, it gets a timestamp. Timestamp is also used to resolve any conflict between critical section requests.
- All algorithm which follows non-token based approach maintains a logical clock. Logical clocks get updated according to Lamport's scheme
- Example: Lamport's algorithm, Ricart–Agrawala algorithm

Quorum based approach:

- Instead of requesting permission to execute the critical section from all other sites, Each site requests only a subset of sites which is called a quorum.
- Any two subsets of sites or Quorum contains a common site.
- This common site is responsible to ensure mutual exclusion
- Example: Maekawa's Algorithm

Deadlock Handling:

The following are the strategies used for Deadlock Handling in Distributed System:

1. **Deadlock Prevention:** As the name implies, this strategy ensures that deadlock can never happen because system designing is carried out in such a way. If any one of the deadlock-causing conditions is not met then deadlock can be prevented. Following are the three

methods used for preventing deadlocks by making one of the deadlock conditions to be unsatisfied:

Collective Requests: In this strategy, all the processes will declare the required resources for their execution beforehand and will be allowed to execute only if there is the availability of all the required resources. When the process ends up with processing then only resources will be released. Hence, the hold and wait condition of deadlock will be prevented.

But the issue is initial resource requirements of a process before it starts are based on an assumption and not because they will be required. So, resources will be unnecessarily occupied by a process and prior allocation of resources also affects potential concurrency.

Ordered Requests: In this strategy, ordering is imposed on the resources and thus, process requests for resources in increasing order. Hence, the circular wait condition of deadlock can be prevented.

An ordering strictly indicates that a process never asks for a low resource while holding a high one.

There are two more ways of dealing with global timing and transactions in distributed systems, both of which are based on the principle of assigning a global timestamp to each transaction as soon as it begins.

It is better to give priority to the old processes because of their long existence and might be holding more resources.

It also eliminates starvation issues as the younger transaction will eventually be out of the system.

Pre-emption: Resource allocation strategies that reject no-pre-emption conditions can be used to avoid deadlocks.

Wait-die: If an older process requires a resource held by a younger process, the latter will have to wait. A young process will be destroyed if it requests a resource controlled by an older process.

Wound-wait: If an old process seeks a resource held by a young process, the young process will be pre-empted, wounded, and killed, and the old process will resume and wait. If a young process needs a resource held by an older process, it will have to wait.

2. Deadlock Avoidance: In this strategy, deadlock can be avoided by examining the state of the system at every step. The distributed system reviews the allocation of resources and wherever it finds an unsafe state, the system backtracks one step and again comes to the

safe state. For this, resource allocation takes time whenever requested by a process. Firstly, the system analysis occurs whether the granting of resources will make the system in a safe state or unsafe state then only allocation will be made.

A safe state refers to the state when the system is not in deadlocked state and order is there for the process regarding the granting of requests.

An unsafe state refers to the state when no safe sequence exists for the system. Safe sequence implies the ordering of a process in such a way that all the processes run to completion in a safe state.

3. Deadlock Detection and Recovery: In this strategy, deadlock is detected and an attempt is made to resolve the deadlock state of the system. These approaches rely on a Wait-For-Graph (WFG), which is generated and evaluated for cycles in some methods.

The following two requirements must be met by a deadlock detection algorithm:

Progress: In a given period, the algorithm must find all existing deadlocks. There should be no deadlock existing in the system which is undetected under this condition. To put it another way, after all, wait-for dependencies for a deadlock have arisen, the algorithm should not wait for any additional events to detect the deadlock.

No False Deadlocks: Deadlocks that do not exist should not be reported by the algorithm which is called phantom or false deadlocks.

There are different types of deadlock detection techniques:

Centralized Deadlock Detector: The resource graph for the entire system is managed by a central coordinator. When the coordinator detects a cycle, it terminates one of the processes involved in the cycle to break the deadlock. Messages must be passed when updating the coordinator's graph. Following are the methods:

A message must be provided to the coordinator whenever an arc is created or removed from the resource graph.

Hierarchical Deadlock Detector: In this approach, deadlock detectors are arranged in a hierarchy. Here, only those deadlocks can be detected that fall within their range.

Distributed Deadlock Detector: In this approach, detectors are distributed so that all the sites can fully participate to resolve the deadlock state. In one of the following below four classes for the Distributed Detection Algorithm- The probe-based scheme can be used for this purpose. It follows local WFGs to detect local deadlocks and probe messages to detect global deadlocks.

There are four classes for the Distributed Detection Algorithm:

- *Path-pushing*: In path-pushing algorithms, the detection of distributed deadlocks is carried out by maintaining an explicit global WFG.
- *Edge-chasing*: In an edge-chasing algorithm, probe messages are used to detect the presence of a cycle in a distributed graph structure along the edges of the graph.
- *Diffusion computation*: Here, the computation for deadlock detection is dispersed throughout the system's WFG.
- *Global state detection*: The detection of Distributed deadlocks can be made by taking a snapshot of the system and then inspecting it for signs of a deadlock.

Implementation:

Suzuki-Kasami Algorithm:

1. To enter Critical section:

- When a site S_i wants to enter the critical section and it does not have the token then it increments its sequence number $RN_i[i]$ and sends a request message **REQUEST(i, sn)** to all other sites in order to request the token. Here **sn** is update value of $RN_i[i]$
- When a site S_j receives the request message **REQUEST(i, sn)** from site S_i , it sets $RN_j[i]$ to maximum of $RN_j[i]$ and **sn** i.e $RN_j[i] = \max(RN_j[i], sn)$.
- After updating $RN_j[i]$, Site S_j sends the token to site S_i if it has token and $RN_j[i] = LN[j] + 1$

2. To execute the critical section:

- Site S_i executes the critical section if it has acquired the token.

3. To release the critical section:

After finishing the execution Site S_i exits the critical section and does following:

- sets $LN[i] = RN_i[i]$ to indicate that its critical section request $RN_i[i]$ has been executed
- For every site S_j , whose ID is not present in the token queue **Q**, it appends its ID to **Q** if $RN_i[j] = LN[j] + 1$ to indicate that site S_j has an outstanding request.
- After above updation, if the Queue **Q** is non-empty, it pops a site ID from the **Q** and sends the token to site indicated by popped ID.
- If the queue **Q** is empty, it keeps the token

Message Complexity:

The algorithm requires 0 message invocation if the site already holds the idle token at the time of critical section request or maximum of N message per critical section execution.

This N messages involves

- (N – 1) request messages
- 1 reply message

Code:

```
import keyboard
import time
import threading

runningP = -1

# RN arrays of processes
RN = {
    0: [0, 0, 0, 0, 0],
    1: [0, 0, 0, 0, 0],
    2: [0, 0, 0, 0, 0],
    3: [0, 0, 0, 0, 0],
    4: [0, 0, 0, 0, 0]
}

token = {
    "token_owner": 2,
    "Q": [],
    "LN": [0, 0, 0, 0, 0],
    "isRunning": False
}

def dispCurrentRNState():
    for key, value in RN.items():
        print(key, ":", value)

def updateRN(processNo, sequenceNumber):
    for key, value in RN.items():
        value[processNo] = max(value[processNo], sequenceNumber)

# Execute cs and remaining tasks

def executeCS(processForCS):
    print("\n*****\n")
    print(f"Process {processForCS} executing CS...")
    print('Token owner is: {}'.format(token["token_owner"]))

    time.sleep(10)
    print(f'\nProcess {processForCS} has completed running CS')

    # Process completed CS
    token["isRunning"] = False

    #update LN
    token["LN"][processForCS] = RN[processForCS][processForCS]
    # print(f"Process Completed CS")

#Check For Outstanding Requests
```

```

    # For every site Sj, whose ID is not present in the token queue Q, it appends
    its ID to Q if  $RN_i[j] = LN[j] + 1$  to indicate that site Sj has an outstanding
    request.
    for index, val in enumerate(RN[token["token_owner"]]):
        # print("Running P: ", runningP)
        if(val == token["LN"][index] + 1 and index != runningP and index not in
        token["Q"]):
            # outstanding Requests
            print(f'Process {index}\''s request is outstanding, it will be added to
            Token\'s Queue')
            token["Q"].append(index)
            print(f'Queue: {token["Q"]}')

    #Handing out the token
    if(len(token["Q"]) != 0):
        # pop a process from the queue and give it the token
        poppedPs = token["Q"].pop(0)
        token["token_owner"] = poppedPs
        token["isRunning"] = True
        executeCS(poppedPs)

if __name__ == "__main__":
    # print("Press Key E to exit")
    print("Running Main Again")

    # Display Current State of RN Arrays
    print("Current RN Arrays: ")
    dispCurrentRNState()
    print(" ")
    print('Token owner is: {}'.format(token["token_owner"]))

    while True:

        if(token["isRunning"]):

            processes = input(
                "Enter Process Numbers which want to access C.S separated
                by space (Click N for None): ")

            if(processes != 'N'):
                psList = processes.strip().split(" ")
                print(" ")

                for ps in psList:
                    processForCS = int(ps)
                    print(f"***** Process {processForCS} *****")
                    seqNo = RN[processForCS][processForCS]+1
                    # Broadcasting Request

                    print(f"Process No.: {processForCS}")
                    print(f"Sequence No.: {seqNo}")

```

```

        print(f"Broadcasting Request ({processForCS} ,
{seqNo}) .....")

        time.sleep(2)
        print("Broadcast complete")
        print(" ")

        # Updating RN Arrays
        print("Updating RN Arrays at all process sites")
        updateRN(processForCS, seqNo)
        print("Current RN Arrays: ")
        dispCurrentRNState()
        print(" ")

    else:

        processForCS = int(input("Enter Process No. which wants to
access C.S: "))
        seqNo = RN[processForCS][processForCS]+1
        # Broadcasting Request

        print(f"Process No.: {processForCS}")
        print(f"Sequence No.: {seqNo}")
        print(f"Broadcasting Request ({processForCS} , {seqNo})
.....")

        time.sleep(2)
        print("Broadcast complete")
        print(" ")

        # Updating RN Arrays
        print("Updating RN Arrays at all process sites")
        updateRN(processForCS, seqNo)
        print("Current RN Arrays: ")
        dispCurrentRNState()
        print(" ")

        # Check condition of sending token: RNj[i] = LN[i] + 1
        if(RN[token["token_owner"]][processForCS] ==
token["LN"][processForCS] + 1):
            # give the token
            print(f"Conditions met, giving token to
{processForCS}...")

            token["token_owner"] = processForCS

            # print(f"New Token Owner: {token["token_owner"]}")
            print('Token owner is: {}'.format(token["token_owner"]))

            token["isRunning"] = True
            runningP = processForCS

            thread = threading.Thread(target=executeCS,
args=(processForCS, ))
            thread.start()

```



```
print("Main Continuing Running")

if keyboard.is_pressed('E'):
    break
```

Output:

Current RN Arrays:

0 : [0, 0, 0, 0, 0]

1 : [0, 0, 0, 0, 0]

2 : [0, 0, 0, 0, 0]

3 : [0, 0, 0, 0, 0]

4 : [0, 0, 0, 0, 0]

Token owner is: 2

Enter Process No. which wants to access C.S: 1

Process No.: 1

Sequence No.: 1

Broadcasting Request (1 , 1)

Broadcast complete

Updating RN Arrays at all process sites

Current RN Arrays:

0 : [0, 1, 0, 0, 0]

1 : [0, 1, 0, 0, 0]

2 : [0, 1, 0, 0, 0]

3 : [0, 1, 0, 0, 0]

4 : [0, 1, 0, 0, 0]

Conditions met, giving token to 1...

Token owner is: 1

Main Continuing Running

Process 1 executing CS...

Token owner is: 1

Enter Process Numbers which want to access C.S separated by space (Click N for None): 2 3

***** Process 2 *****

Process No.: 2

Sequence No.: 1

Broadcasting Request (2 , 1)

Broadcast complete

Updating RN Arrays at all process sites

Current RN Arrays:

0 : [0, 1, 1, 0, 0]

1 : [0, 1, 1, 0, 0]

2 : [0, 1, 1, 0, 0]

3 : [0, 1, 1, 0, 0]

4 : [0, 1, 1, 0, 0]

***** Process 3 *****

Process No.: 3

Sequence No.: 1

Broadcasting Request (3 , 1)

Broadcast complete

Updating RN Arrays at all process sites

Current RN Arrays:

0 : [0, 1, 1, 1, 0]

1 : [0, 1, 1, 1, 0]

2 : [0, 1, 1, 1, 0]

3 : [0, 1, 1, 1, 0]

4 : [0, 1, 1, 1, 0]

Enter Process Numbers which want to access C.S separated by space (Click N for None):

Process 1 has completed running CS

Process 2's request is outstanding, it will be added to Token's Queue

Queue: [2]

Process 3's request is outstanding, it will be added to Token's Queue

Queue: [2, 3]

Process 2 executing CS...

Token owner is: 2

Process 2 has completed running CS

Process 3 executing CS...

Token owner is: 3

Process 3 has completed running CS

Link to Code File:

https://drive.google.com/file/d/1HIm9UgHaVARTs_4rNj8FC6nQgKMKpgnJ/view?usp=share_link

Conclusion:

In conclusion, our experiment focused on studying mutual exclusion and deadlock handling in a distributed system. We implemented a simulation of the Suzuki-Kasami algorithm as a solution to the challenges we identified. The algorithm demonstrated how processes in a distributed system can coordinate with each other to prevent race conditions and ensure mutual exclusion, without causing deadlocks.

Postlab Questions:

1. Explain the different ways of recovery from deadlock.
2. What are the features of CMH algorithm